

Cosine Trends:

- The seasonal means model is not very efficient if the seasonal pattern is smooth
 - We can define a function that might capture the seasonality using fewer coefficients using a periodic function
- It needs to meet at the end of one cycle and beginning of the next

$$\mu_t = \beta \cos(2\pi f t + \Phi)$$

- where β is the amplitude,
- f is the frequency of the oscillation (assumed known for seasonal ts)
- and Φ is the phase of the curve

- The curve repeats itself every $1/f$ time units = PERIOD of the curve
 - $f=1/12 \Rightarrow \text{Period}=12$
- To estimate β and Φ , we would need to have access to nonlinear regression techniques.
 - Nonlinear regression models can be generally written as:
- Nonlinear least squares defines the same target as OLS: $(y-f(x))^2$
 - Nonlinear regression models are distinguished from linear models by the derivatives with respect to the parameters (slopes) are functions of the slopes
 - Consider the partial derivatives of a linear model with $\log(x_1)$, x_2 , and $\log(y)$:
- NLS uses an iterative weighted least squares algorithm to find the coefficient estimates and an approximate variance-covariance matrix of the coefficients
 - For a brief discussion of `nlS` see: <http://cran.r-project.org/doc/contrib/Fox-Companion/appendix-nonlinear-regression.pdf>
 - The main differences in using `nlS` vs `lm` involve writing out all the model components and providing initial estimates of the parameters.
 - An `nlS` version of an SLR model is:
- For the cosine seasonal trend model using NLS:
 - CC assumes that f is known initially
 - This corresponds to our definition of seasonality but not cycles
 - As long as we are careful about the parameter initialization, we can consider the nonlinear version using `nlS`
 - later we could consider using `gnlS` from the `nlme` package
- CC uses a trig identity to re-express (re-parametrize) the previous nonlinear regression model into the following linear model:

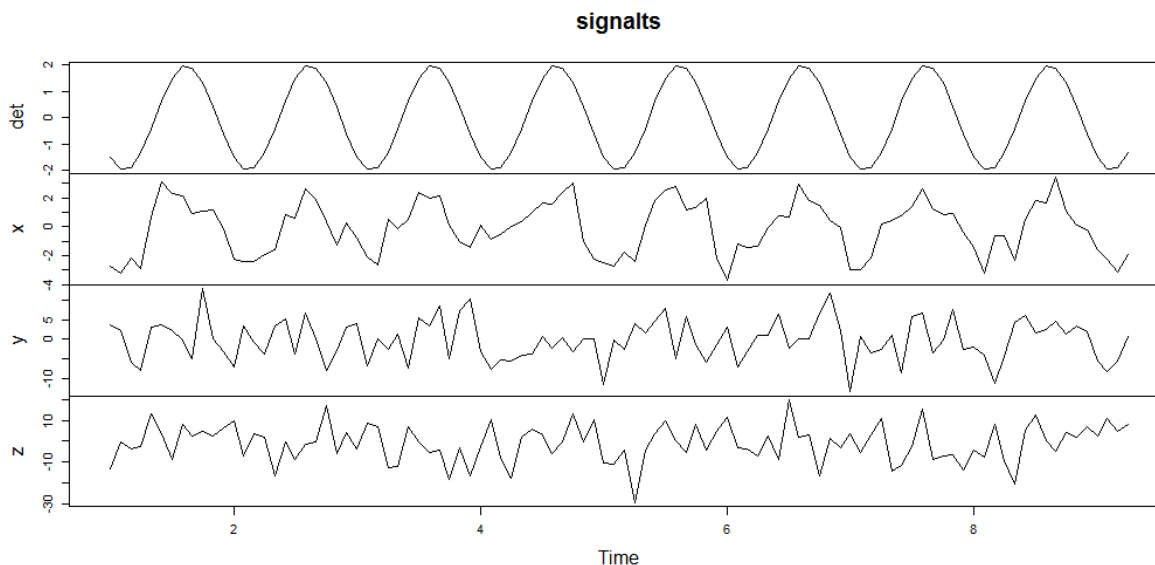
$$\beta \cos(2\pi f t + \Phi) =$$

- We really would want to work with the following trend model that incorporates a constant mean (or a cosine at frequency 0):

$$\mu_t = \beta_0 + \beta \cos(2\pi f t + \Phi) =$$

- To generate the sine and cosines, the TSA package contains the function `harmonic`
- Also note the very useful advice on p. 34 about how you code “time” and how that coding can change the definition of frequency
- If we continued this process for all the different frequencies then we would decompose the original series into a combination of sines and cosines at different frequencies, which is exactly what a fourier analysis actually does.

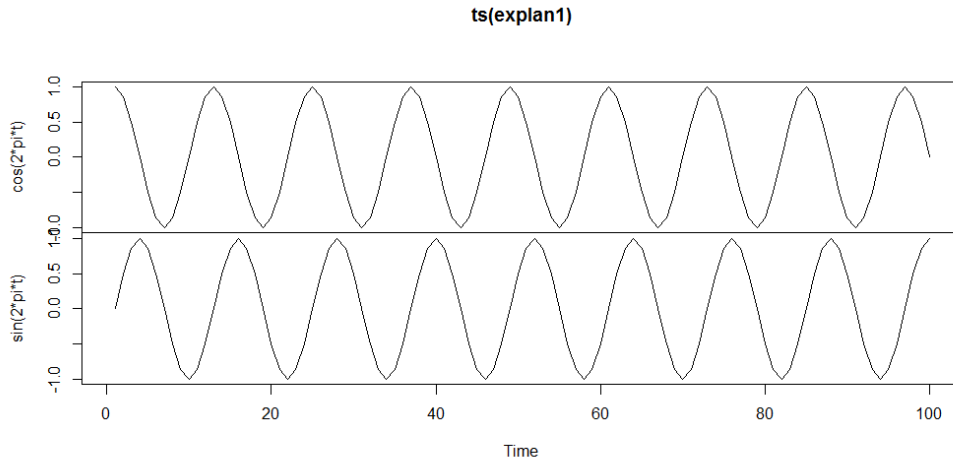
```
> time<-seq(1: 100)
> det<-2*cos(2*pi *time/12+. 6*pi )
> set.seed(13579)
> x<-det+rnorm(100, 0, 1)
> y<-det+rnorm(100, 0, 5)
> z<-det+rnorm(100, 0, 10)
> signal ts<-ts(data.frame(det, x, y, z), freq=12)
> plot.ts(signal ts)
```



```
> require(TSA)
> expl an1<-harmonic(signal ts, m=1)
> head(data.frame(Month=season(signal ts), round(expl an1, 2)), 16)
```

	Month	cos. 2. pi . t.	sin. 2. pi . t.
1	January	1.00	0.00
2	February	0.87	0.50
3	March	0.50	0.87
4	April	0.00	1.00
5	May	-0.50	0.87
6	June	-0.87	0.50
7	July	-1.00	0.00
8	August	-0.87	-0.50
9	September	-0.50	-0.87
10	October	0.00	-1.00
11	November	0.50	-0.87
12	December	0.87	-0.50
13	January	1.00	0.00
14	February	0.87	0.50
15	March	0.50	0.87
16	April	0.00	1.00

```
> plot(ts(explan1))
```



```
> l m1<-l m(det~harmoni c(si gnal ts), data=si gnal ts)
```

```
> summary(l m1)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.887e-16	6.312e-16	2.990e-01	0.766
harmoni c(si gnal ts)cos(2*pi *t)	-1.486e+00	8.923e-16	-1.666e+15	<2e-16 ***
harmoni c(si gnal ts)si n(2*pi *t)	-1.338e+00	8.923e-16	-1.500e+15	<2e-16 ***

Residual standard error: 6.305e-15 on 97 degrees of freedom
Multiple R-squared: 1, Adjusted R-squared: 1
F-statistic: 2.553e+30 on 2 and 97 DF, p-value: < 2.2e-16

```
> anova(l m1)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
harmoni c(si gnal ts)	2	203	101.5	2.5529e+30	< 2.2e-16 ***
Resi dual s	97	0	0.0		

Warning message:

In anova.lm(l m1) :

ANOVA F-tests on an essentially perfect fit are unreliable

```
> fi t1<-fi tted(l m1)
```

```
> sqrt(sum(l m1$coef[2:3]^2))
```

```
[1] 2
```

```
>
```

```
> l m2<-l m(x~harmoni c(si gnal ts), data=si gnal ts)
```

```
> summary(l m2)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.01511	0.09245	-0.163	0.871
harmoni c(si gnal ts)cos(2*pi *t)	-1.83420	0.13069	-14.034	< 2e-16 ***
harmoni c(si gnal ts)si n(2*pi *t)	-1.21600	0.13069	-9.304	4.27e-15 ***

Residual standard error: 0.9235 on 97 degrees of freedom
Multiple R-squared: 0.7479, Adjusted R-squared: 0.7428
F-statistic: 143.9 on 2 and 97 DF, p-value: < 2.2e-16

```
> anova(l m2)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
harmoni c(si gnal ts)	2	245.488	122.744	143.92	< 2.2e-16 ***
Resi dual s	97	82.727	0.853		

```
> fi t2<-fi tted(l m2)
```

```
> sqrt(sum(l m2$coef[2:3]^2))
```

```
[1] 2.200664
```

```
>
```

```
> l m3<-l m(y~harmoni c(si gnal ts), data=si gnal ts)
```

```
> summary(l m3)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.1441	0.4845	-0.297	0.76680
harmoni c(si gnal ts)cos(2*pi *t)	-1.9561	0.6849	-2.856	0.00525 **
harmoni c(si gnal ts)si n(2*pi *t)	-1.8924	0.6849	-2.763	0.00685 **

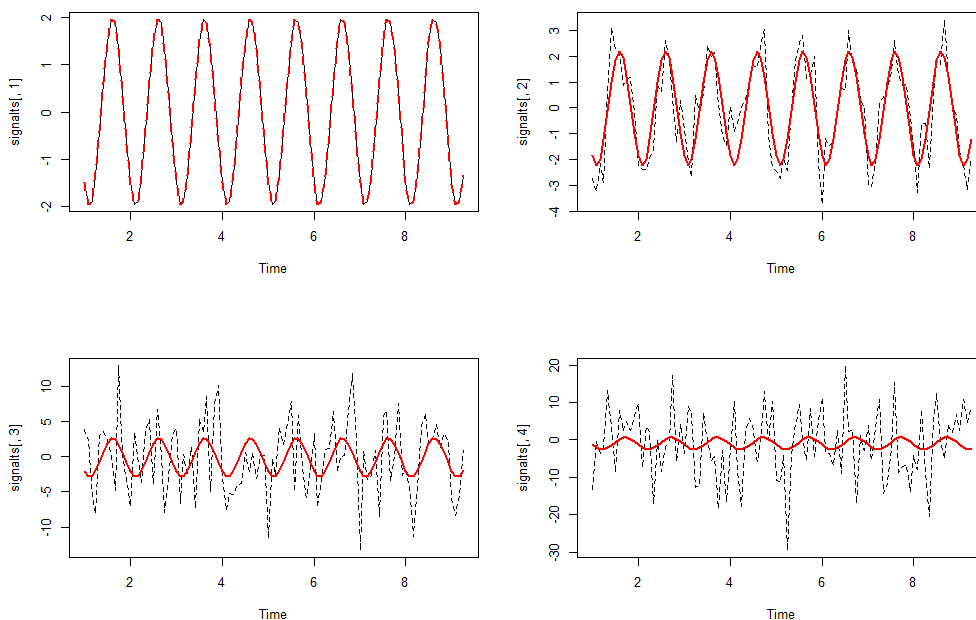
Residual standard error: 4.839 on 97 degrees of freedom
Multiple R-squared: 0.142, Adjusted R-squared: 0.1243
F-statistic: 8.027 on 2 and 97 DF, p-value: 0.0005946

```
> fi t3<-fi tted(lm3)
> lm4<-lm(z-harmonic(signal ts), data=signal ts)
> summary(lm4)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.9643	0.9041	-1.067	0.289
harmonic(signal ts)cos(2*pi*t)	-0.3708	1.2781	-0.290	0.772
harmonic(signal ts)sin(2*pi*t)	-1.5176	1.2781	-1.187	0.238

Residual standard error: 9.031 on 97 degrees of freedom
Multiple R-squared: 0.01528, Adjusted R-squared: -0.00502
F-statistic: 0.7528 on 2 and 97 DF, p-value: 0.4738

```
> fi t4<-fi tted(lm4)
> par(mfrow=c(2, 2))
> plot(signal ts[, 1])
> lines(fi t1~as.vector(time(signal ts)), col="red", lty=2, lwd=2)
> plot(signal ts[, 2], lty=2)
> lines(fi t2~as.vector(time(signal ts)), col="red", lty=1, lwd=2)
> plot(signal ts[, 3], lty=2)
> lines(fi t3~as.vector(time(signal ts)), col="red", lty=1, lwd=2)
> plot(signal ts[, 4], lty=2)
> lines(fi t4~as.vector(time(signal ts)), col="red", lty=1, lwd=2)
```



Or we can fit the models more directly using nls:

```
> 0.6*pi
[1] 1.884956
>
> model 1<-nls(signal ts[, 1]~a+b*cos(2*pi*time/12+c), start=list(a=0, b=2, c=0.6*pi))
Error in nls(signal ts[, 1] ~ a + b * cos(2 * pi * time/12 + c), start = list(a = 0, b = 2, c = 0.6 * pi)) :
  number of iterations exceeded maximum of 50
>
> model 2<-nls(signal ts[, 2]~a+b*cos(2*pi*time/12+c), start=list(a=0, b=2, c=0.6*pi))
> summary(model 2)
```

Formula: signal ts[, 2] ~ a + b * cos(2 * pi * time/12 + c)

Parameters:

	Estimate	Std. Error	t value	Pr(> t)
a	-0.01511	0.09245	-0.163	0.871
b	2.20066	0.12971	16.966	<2e-16 ***
c	2.03256	0.05983	33.972	<2e-16 ***

```

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9235 on 97 degrees of freedom

Number of iterations to convergence: 3
Achieved convergence tolerance: 2.132e-08

> model 3<-nl s(signal ts[, 3]~a+b*cos(2*pi *time/12+c), start=list(a=0, b=2, c=0.6*pi))
> summary(model 3)

Formula: signal ts[, 3] ~ a + b * cos(2 * pi * time/12 + c)

Parameters:
      Estimate Std. Error t value Pr(>|t|)
a    -0.1441     0.4845  -0.297 0.766795
b     2.7217     0.6793   4.007 0.000121 ***
c     1.8491     0.2537   7.290 8.39e-11 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.839 on 97 degrees of freedom

Number of iterations to convergence: 3
Achieved convergence tolerance: 9.352e-09

>
> model 4<-nl s(signal ts[, 4]~a+b*cos(2*pi *time/12+c), start=list(a=0, b=2, c=0.6*pi))
> summary(model 4)

Formula: signal ts[, 4] ~ a + b * cos(2 * pi * time/12 + c)

Parameters:
      Estimate Std. Error t value Pr(>|t|)
a    -0.9643     0.9041  -1.067  0.289
b     1.5622     1.2733   1.227  0.223
c     1.2869     0.8212   1.567  0.120

Residual standard error: 9.031 on 97 degrees of freedom

Number of iterations to convergence: 4
Achieved convergence tolerance: 2.967e-08

Note that although one method used nls and the other ls, that the models are actually equivalent here (same likelihood and # parameters):
> AIC(lm4, model 4)
      df      AIC
lm4      4 728.8845
model 4  4 728.8845

> mean(abs(fitted(lm4)-fitted(model 4)))
[1] 2.270966e-07

```

- Note that this equivalence may not always hold due to differences in numerical maximization routines and parameter initializations between `nl s` and linearized versions of the models.
 - It does suggest that although the initial version of the model is nonlinear it is actually close to the linearized version of the model
 - The inference in the nonlinear model is probably pretty safe
 - Stability of `nl s` is related to how linear the model is, with a and b entering the nonlinear model linearly, the only coefficient with “nonlinearity” attached to it is the phase coefficient, c .